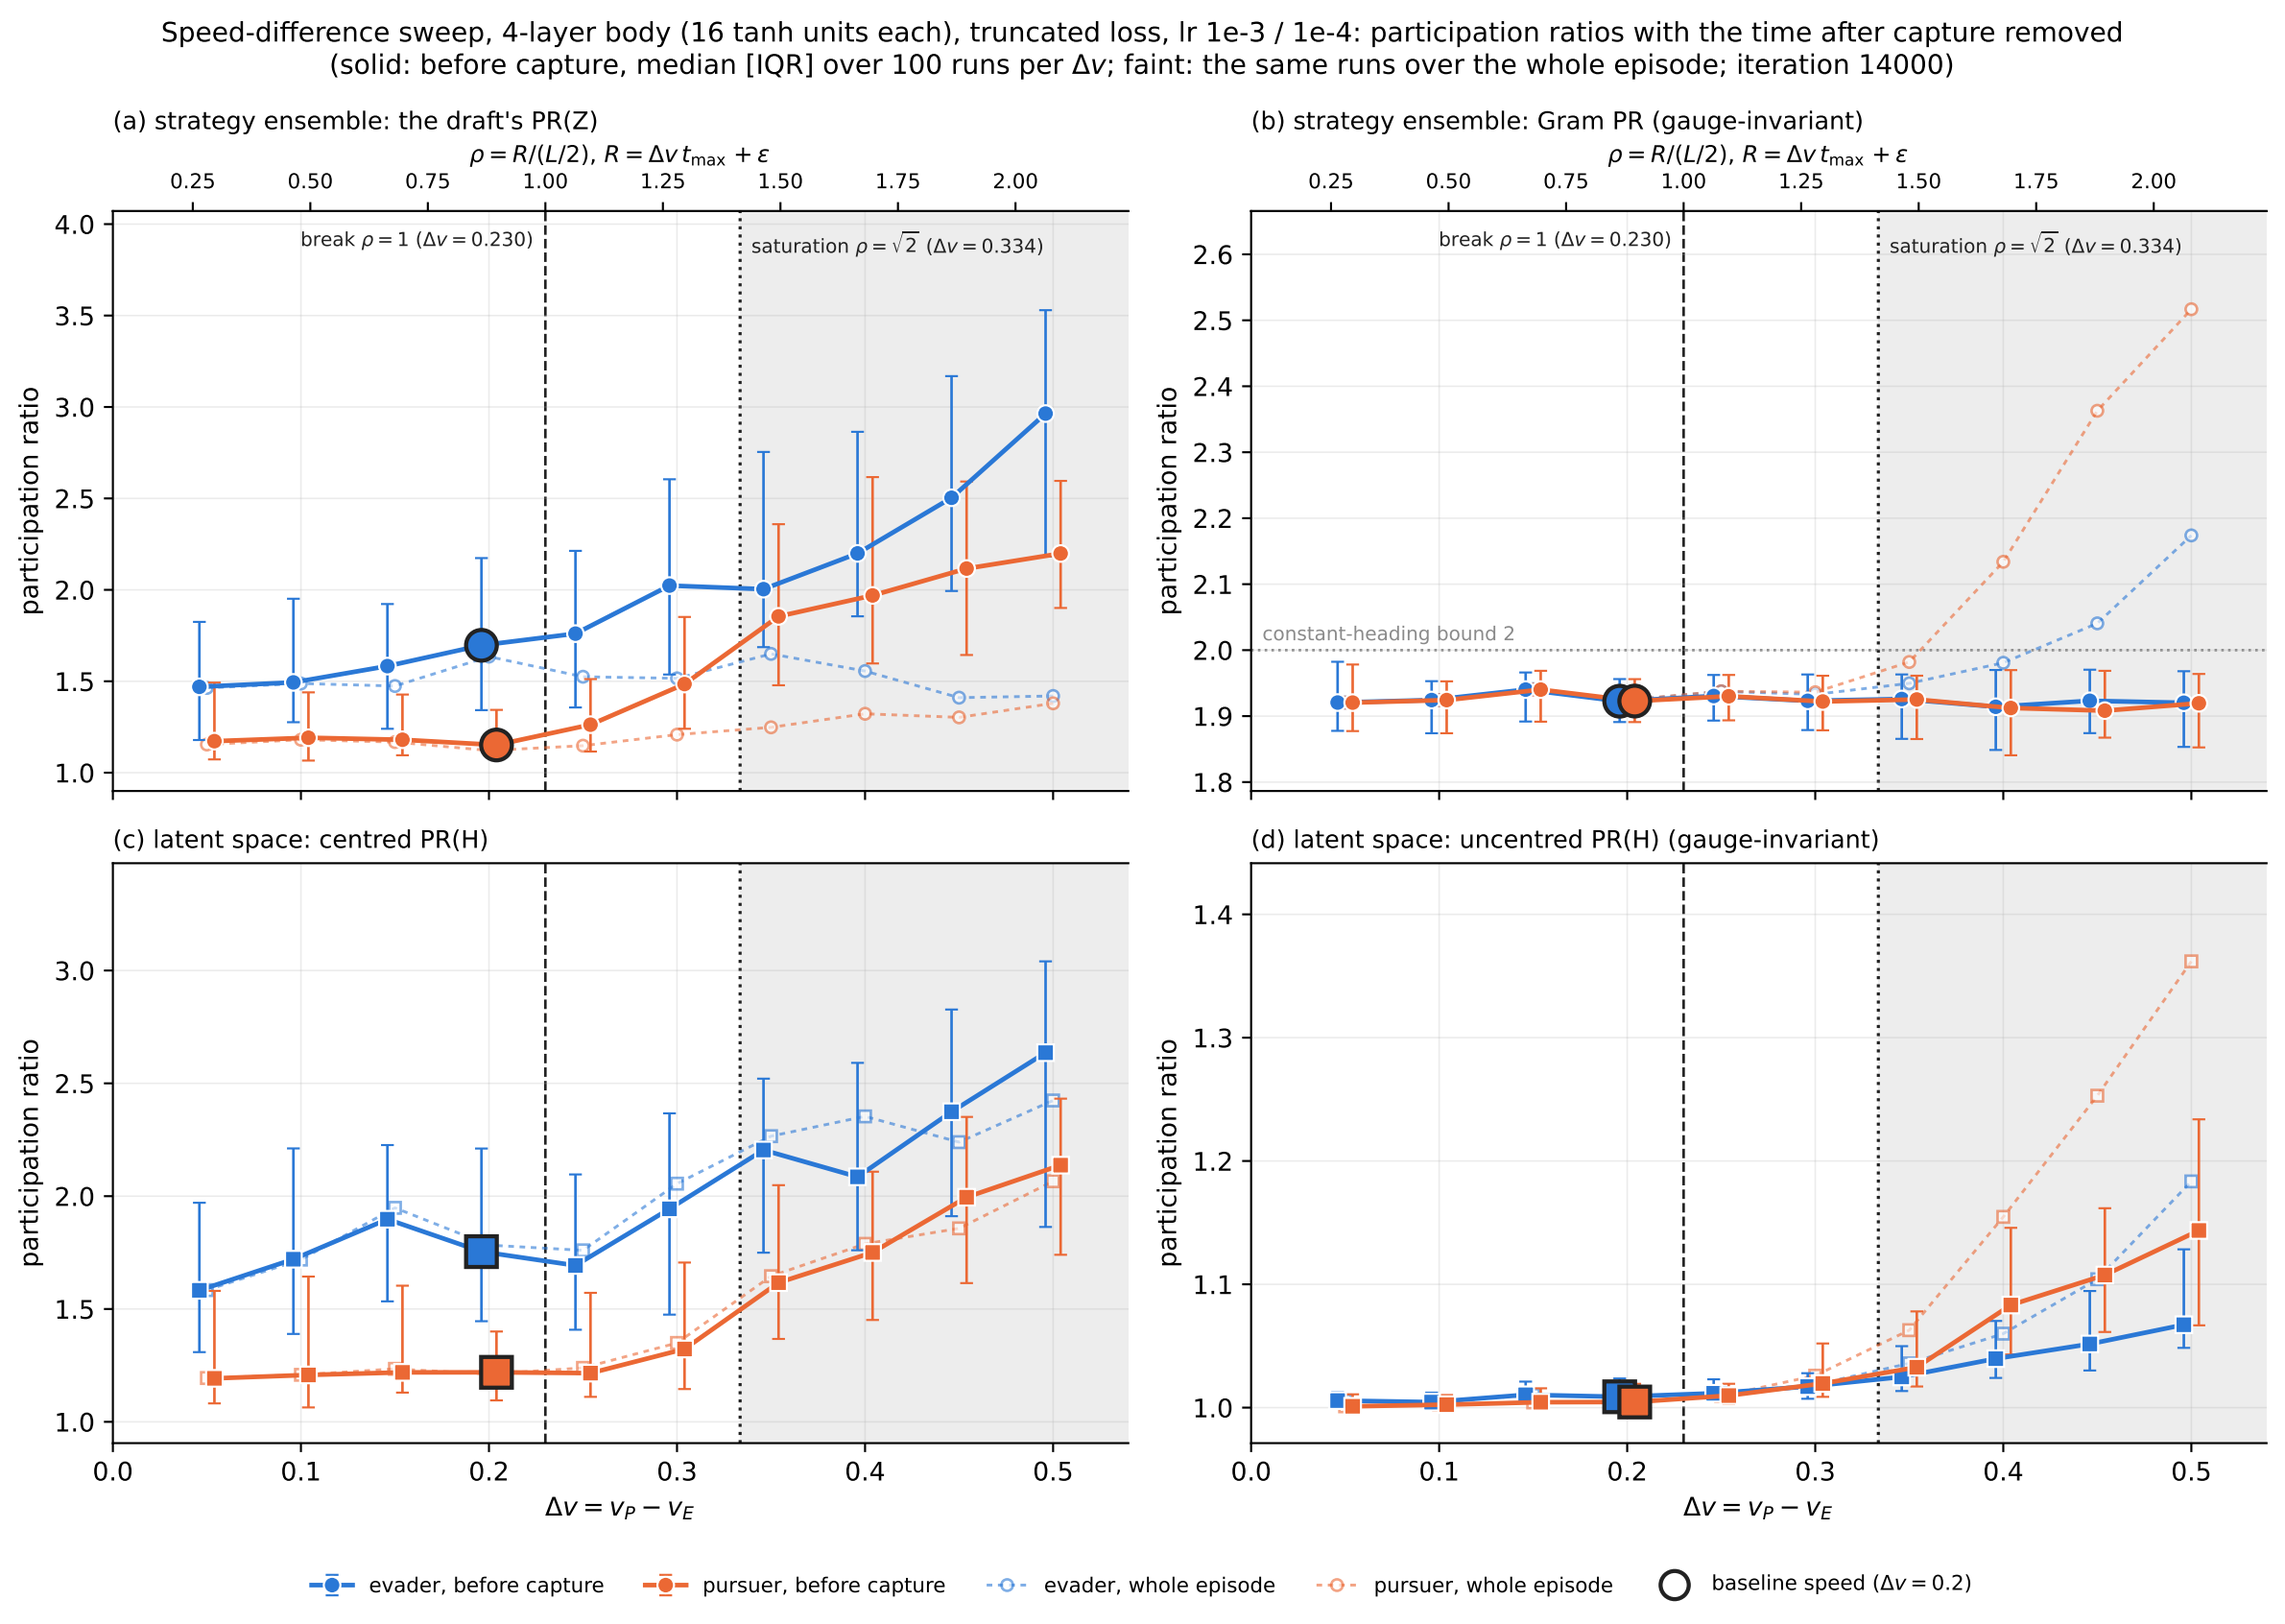
Speed-difference sweep, 4-layer body (16 tanh units each), truncated loss, lr 1e-3 / 1e-4: participation ratios with the time after capture removed
(solid: before capture, median [IQR] over 100 runs per Δv; faint: the same runs over the whole episode; iteration 14000)
(a) strategy ensemble: the draft's PR(Z)
(b) strategy ensemble: Gram PR (gauge-invariant)
(c) latent space: centred PR(H)
(d) latent space: uncentred PR(H) (gauge-invariant)
ρ = R/(L/2), R = Δv t_max + ε
break ρ = 1 (Δv = 0.230)
saturation ρ = √2 (Δv = 0.334)
constant-heading bound 2
participation ratio
Δv = v_P − v_E
evader, before capture
pursuer, before capture
evader, whole episode
pursuer, whole episode
baseline speed (Δv = 0.2)